

Rational Heuristics for One-Shot Games*

Fredrick Callaway¹, Thomas L. Griffiths², and Gustav Karreskog^{†3}

^{1,2}Department of Psychology, Princeton

³Department of Economics, Stockholm School of Economics

November 11, 2020

JOB MARKET PAPER

For the most recent version, [click here](#)

Abstract

We present a theory of human behavior in one-shot games as being the result of rational use of heuristics. To test this theory, we define a broad family of heuristics for one-shot games and an associated family of cognitive cost functions. In a large, preregistered experiment, we find that behavior is quantitatively well predicted by heuristic reasoning and optimal use of limited cognitive resources, and yields better predictions than existing models based on iterated reasoning. Furthermore, we find that, holding the game constant, the participants' actions depend on their previous experience in the way predicted by the rational use of heuristics. This theory has implications for the broader field of behavioral economics and can help us understand when and why different deviations from the rational benchmark appear.

Keywords: Bounded rationality, Experiments, Cognitive cost, Strategic thinking, Game theory, One-Shot games, Heuristics

JEL classification: C72, C90, D83, D01

*We thank Drew Fudenberg, Alice Hallman, Benjamin Mandl, Erik Mohlin, Isak Trygg Kuper-smidt, Jörgen Weibull, Peter Wikman, and seminar participants at SSE, SUDSWEC, UCL, NHH, and Princeton, for helpful comments and insights. This work was supported by a grant to TLG by the Templeton foundation, Tom Hedelius and Knut and Alice Wallenberg Research Foundation,

[†]gustav.karreskog@phdstudent.hhs.se

1 Introduction

A key assumption underlying classical economic theory is that people behave optimally, in order to maximize their expected utility (Savage, 1954). However, a large body of work in behavioral economics show that human behavior systematically deviates from this rational benchmark (Kahneman, 2011) in many settings. This suggest we can improve our understanding by incorporating more realistic and behavioral components to our models of economic behavior. While many of these deviations are indeed systematic and show up in multiple studies, the estimated biases vary considerably between studies and contexts. Apparent biases change considerably or even disappear if participants have opportunities for learning or if the details of the decision task changes. For example, this is the case for the endowment effect (Tunçel and Hammitt, 2014), loss aversion (Ert and Erev, 2013), and present bias (Imai, Rutter and Camerer, 2020).

In order to incorporate behavioral effects into theories with broader applications, without having to run new experiments for every specific setting, we need a theory that can account for these variations. That is, we need a theory that can help us understand why, and predict when, we should expect deviations from the rational benchmark, and when we can safely assume behavior is close to rational. In this paper, we propose such theory where people use simple decision procedures, or *heuristics*, that are optimized to the environment in order to make the best possible use of their limited cognitive resources and thereby maximize utility. This allows us to predict behavior by analyzing which heuristics perform well in which environments. In this paper, we present an explicit version of this theory tailored to one-shot games and test it experimentally.

In situations where people play the same game multiple times against different opponents, so that there is an opportunity for learning, both theoretical and experimental work suggests that Nash Equilibrium can be a sensible long-run prediction in many cases (Fudenberg et al., 1998; Camerer, 2003). However, in experimental studies of one-shot games where players don't have experience of the particular game at hand, people seldom follow the theoretical prediction of Nash equilibrium play (see Crawford, Costa-Gomes and Iriberri, 2013 for an overview). Consequently, we need an alternative theory for strategic interactions that only happen once (or infrequently).

The most common theories for behavior in one-shot games in the literature assume that players perform some kind of iterated reasoning to form beliefs about the other player's action, and then select the best action in response. Examples of such models are so-called level-k models, introduced by Nagel (1995); Stahl and Wilson (1994, 1995), and closely related Cognitive Hierarchy (CH) models, introduced by Camerer, Ho and Chong (2004). In such models, participants are characterized by different levels of reasoning. Level-0 reasoners behave naively, often assumed to play a uniformly random strategy. Level-1 reasoners best respond to level-0 behavior, and even higher levels best respond to behavior based on lower level reasoning. In meta-analyses such as Crawford, Costa-Gomes and Iriberri (2013), Wright and Leyton-Brown (2017), and Fudenberg and Liang (2019), variations of these iterated reasoning models best explain human behavior.

A key assumption underlying iterated reasoning models is the basic structure of belief formation and best response. However, empirical evidence on information acquisition and elicited beliefs is often inconsistent with a such a belief-response process. When participants are asked to state their beliefs about how the opponent will play, they often fail to play a best response to those beliefs (Costa-Gomes and Weizsäcker, 2008). Eye-tracking studies have revealed that the order in which participants attend to payoffs in visually presented normal-form games is inconsistent with a belief-formation and best-response process (Polonio, Di Guida and Coricelli, 2015; Devetag, Di Guida and Polonio, 2016; Stewart et al., 2016). Furthermore, the estimated parameters of these models often vary considerably between different data sets, behavior seems to depend on the underlying game in a way not captured by the models (Bardsley et al., 2010; Heap, Arjona and Sugden, 2014), and there is evidence of earlier games played having an effect on behavior not captured by existing models (Mengel and Scitubba, 2014; Peysakhovich and Rand, 2016).

In this paper, we present a theory of human behavior in one-shot games based on the rational use of heuristics (Lieder and Griffiths, 2017, 2019). That is, we assume that people use simple cognitive strategies that flexibly and selectively process payoff information in order to construct a decision with minimal cognitive effort. These heuristics do not necessarily involve any explicit construction of beliefs to which the players best respond. However, we assume that people adapt the heuristics in order to maximize utility. Although they might not choose the best action in a given game, they will learn which heuristics generally work well in an environment.¹

Our approach thus combines two perspectives on human decision making, embracing both the notion that human behavior is adaptive in a way that can be described as optimization and also the notion that people use simple strategies that are effective for the problems they actually need to solve. The key assumption in this approach, *resource-rational analysis*, is that people use cognitive strategies that make optimal use of their limited computational resources (Lieder and Griffiths, 2019; Griffiths, Lieder and Goodman, 2015 c.f. Lewis, Howes and Singh, 2014; Gershman, Horvitz and Tenenbaum, 2015).

In comparison with traditional rational models, resource-rational analysis is distinctive in that it explicitly accounts for the cost of allocating limited computational resources to a given decision, specifying an objective function that includes both the utility of a decision’s outcome as well as the cost of the cognitive process that produced the decision. In comparison with theories of bounded or ecological rationality (Gigerenzer and Todd, 1999; Goldstein and Gigerenzer, 2002; Smith, 2003; Todd and Gigerenzer, 2012), resource-rational analysis is distinctive in its assumption that people optimize this objective function. This makes it possible to predict when people will use one heuristic versus another (Lieder and Griffiths, 2017) and even to automatically discover novel heuristics (Lieder, Krueger and Griffiths, 2017). Finally, our approach is perhaps most compatible with information-theoretic approaches such as rational inat-

¹This idea is related to that of procedural rationality in Simon (1976).

tention (Matějka and McKay, 2015; Sims, 1998; Caplin and Dean, 2013; Hebert and Woodford, 2019; Steiner, Stewart and Matějka, 2017), in which the costs and benefits of information processing are optimally traded off; resource-rational analysis is distinct however, in making stronger assumptions about the specific computational processes and costs that are likely to be involved in a given domain.

One important commonality between our approach and ecological rationality is the recognition that the quality or adaptiveness of a heuristic depends on the environment in which it is to be used. For example, in an environment in which most interactions are characterized by competing interests (e.g., zero-sum games), a good heuristic is one that looks for actions with high guaranteed payoffs. On the other hand, if most interactions are common interest, focusing on the guaranteed payoff will lead to many missed opportunities for mutually beneficial outcomes, so it might be better to look for the common interest. This is the key insight that allows us to test our theories.

To examine whether people adapt their play as our approach predicts, we conducted a large, preregistered² behavioral experiment. In our experiment, participants play a series of normal form games in one of two environments, which are characterized by different correlations in payoffs. In the *common interest* group, there is a positive correlation between the payoffs of the row and column player over the set of strategy profiles. In the *competing interest* group, the correlation is negative. Interspersed among these treatment-specific games, we include four *comparison games*, which are the same for both conditions (and all sessions). If the participants are using environment-adapted heuristics to make decisions, and different heuristics are good for common interest and competing interest environments, the participants should behave differently in the comparison games since they are employing different heuristics. Indeed, this is what we observe.

To take our analysis further, we define a parameterized family of heuristics and cognitive costs, in order to test the critical resource-rational hypothesis that our participant’s behavior is consistent with an optimal tradeoff between payoff and cognitive cost. Rather than identifying the parameters that best fit human behavior, we identify the parameters that strike this optimal tradeoff, and ask how well they predict the effect of the environment on human behavior. Although we fit the parameters of the cost function that partially defines the resource-rational heuristic—critically—these parameters are fit jointly to data in both treatments. Thus, any differences in predicted behavior is an *a priori* prediction. Strikingly, we find that this model, which has no free parameters that vary between the treatments, achieves nearly the same out-of-sample predictive accuracy as the model with all parameters fit separately to each treatment.

We will start by introducing the general model in Section 2, capturing the connection between heuristics, their associated cognitive costs, the environment, and resource-rationally optimal heuristics. In Section 3 we introduce our main specification of the

²The preregistration is embargoed at the open science foundations preregistration platform. Email the author Gustav Karreskog at gustav.karreskog@phdstudent.hhs.se if you need access to it.

available heuristics and their cognitive costs, metaheuristics. We then introduce the experiment in Section 4, followed by the model-free results based on the comparison games. We there confirm that the two different environments indeed lead to large and predictable differences in behavior. After that we test the two model-based hypothesis using the metaheuristics. Based on these out-of sample predictions we show that the differences in behavior between the different treatments can be accurately predicted by assuming that the participants use the metaheuristics that are optimal in the respective environments. In Section 5 we see that we can confirm the model-based hypothesis also by considering a completely different representation of the possible heuristics using a constrained neural network design, which we call *deep heuristics*. Lastly, in Section 6, we compare our model to a quantal cognitive hierarchy model, and a model with noisy-best reply and pro-social preferences, and show that our model predicts behavior better than both these alternatives.

2 General Model

We consider a setting where individuals in a population are repeatedly randomly matched with another individual to play a finite normal form game. We assume they use some heuristic to decide what strategy to play.

Let $G = \langle \{1, 2\}, S_1 \times S_2, \pi \rangle$ be a two-player normal form game with pure strategy sets $S_i = \{1, \dots, m_i\}$ for $i \in \{1, 2\}$, where $m_i \in \mathbb{N}$. A mixed strategy for player i is denoted $\sigma_i \in \Delta(S_i)$. The *material payoff* for player i from playing pure strategy $s_i \in S_i$ when the other player $-i$ plays strategy $s_{-i} \in S_{-i}$ is denoted $\pi_i(s_i, s_{-i})$. We extend the material payoff function to the expected material payoff from playing a mixed strategy $\sigma_i \in \Delta(S_i)$ against the mixed strategy $\sigma_{-i} \in \Delta(S_{-i})$ with $\pi_i(\sigma_i, \sigma_{-i})$, in the usual way. A heuristic is a function that maps a game to a mixed strategy $h_i(G) \in \Delta(S_i)$. For simplicity, we will always consider the games from the perspective of the row player, and consider the transposed game $G' = \langle \{2, 1\}, S_2 \times S_1, (\pi_2, \pi_1) \rangle$ when talking about the column player's behavior.

Each heuristic has an associated cognitive cost, in general depending on the game as well, $c(h) \in \mathbb{R}_+^3$. Simple heuristics, such as playing the uniformly random mixed strategy, have low cognitive costs, while complicated heuristics involving many precise computations have high cognitive costs. Since a heuristic returns a mixed strategy, the expected material payoff for player i using heuristic h_i when player $-i$ uses heuristic h_{-i} is

$$\pi_i(h_i(G), h_{-i}(G')).$$

Since each heuristic has an associated cognitive cost, the actual expected utility derived

³In general, we can imagine that the cognitive cost depends on both the heuristic and the game, for example it might be more costly to apply it to a 5×5 game than a 2×2 game. But since all our games will be 3, we drop that dependency.

is

$$u_i(h_i, h_{-i}, G) = \pi_i(h_i(G), h_{-i}(G')) - c(h_i).$$

A heuristic is neither good nor bad in isolation; its performance has to be evaluated with regard to some environment, in particular, with regard to the games and other-player behavior one is likely to encounter. Let \mathcal{G} be the set of possible games in the environment, \mathcal{H} be the set of available heuristics, and P be the joint probability distribution over \mathcal{G}, \mathcal{H} . In the equations below we will assume that \mathcal{G} and \mathcal{H} are countable. An environment is given by $\mathcal{E} = (P, \mathcal{G}, \mathcal{H})$. Thus, an environment describes which game and opponent heuristic combinations a player is likely to face. Given an environment we can calculate the expected performance of an heuristic as

$$V(h_i, \mathcal{E}) = \mathbb{E}_{\mathcal{E}} [u_i(h_i, h_{-i}, G)] = \sum_{G \in \mathcal{G}} \sum_{h_{-i} \in \mathcal{H}} u_i(h_i, h_{-i}, G) \cdot P(G, h_{-i}). \quad (1)$$

We can also evaluate the performance of a heuristic conditional on the specific game being played

$$V(h_i, \mathcal{E}, G) = \mathbb{E}_{\mathcal{E}|G} [u_i(h_i, h_{-i}, G)] = \sum_{h_{-i} \in \mathcal{H}} u_i(h_i, h_{-i}, G) \cdot P(h_{-i}|G).$$

We can now define formally what we mean with a rational, or optimal, heuristic. A rational heuristic h^* is an heuristic that optimizes (1), i.e.,

$$h^* = \operatorname{argmax}_{h \in \mathcal{H}} V(h, \mathcal{E}).$$

We here also see that by varying the environment, we can vary which heuristics are optimal. In the experiment, we will vary \mathcal{P} , and thereby vary the predictions we get by assuming rational heuristics.

3 Metaheuristics

To build a formal model of heuristics for one-shot games, we begin by specifying a small set of candidate forms that such a heuristic might take: row-based reasoning, cell-based reasoning, and simulation-based reasoning. For each class, we specify a precise functional form, each employing a small number of continuous parameters, and a cognitive cost function. The cognitive cost of a heuristic is a function of its parameters, and the form of the cost function is itself parameterized. Finally, we consider a higher-order heuristic, which we call a *metaheuristic* that selects among the candidate first-order heuristics based on their expected values for the current game. We emphasize that we do not claim that this specific family captures all the heuristics

people might employ. However, we hypothesized, and our results show, that it is expressive enough to both illustrate the predictions of the general theory, and also provide a strong quantitative explanation of human behavior.

To exemplify the different heuristics, we will apply them to the following example game.

	1	2	3
1	0, 1	0, 2	8, 8
2	5, 6	5, 5	2, 2
3	6, 5	6, 6	1, 1

Figure 1: Example normal form game represented as a bi-matrix. The row player chooses a row and column player chooses a column. The first number in each cell is the payoff of the row player and the second number is the payoff of the column player.

3.1 Row Heuristics

A *row heuristic* calculates a value, $v(s_i)$, for each pure strategy, $s_i \in S_i$, based only on the player's own payoffs associated with s_i . That is, it evaluates a strategy based only on first entries in each cell of the corresponding row of the payoff matrix (see Figure 1). Formally, a row heuristic is defined by the row-value function v such that

$$v(s_i) = f(\pi_i(s_i, 1), \dots, \pi_i(s_i, m_i))$$

for some function $f : \mathbb{R}^m \rightarrow \mathbb{R}$. For example, if f is the mean function, then we have

$$v^{\text{mean}}(s_i) = \frac{1}{m_{-i}} \sum_{s_{-i} \in S_{-i}} \pi_i(s_i, s_{-i}),$$

which results in level-1 like behavior. Indeed, deterministically selecting $\arg \max_{s_i} v^{\text{mean}}(s_i)$ gives exactly the behavior of a level-1 player in the classical level-k model.

If, instead, we let f be min, we recover the maximin heuristic, which calculates the minimum value associated with each strategy and tries to chose the row with highest minimum value,

$$v^{\text{min}}(s_i) = \min_{s_{-i} \in S_{-i}} \pi_i(s_i, s_{-i}),$$

and similarly the maximax heuristic when f is max,

$$v^{\text{max}}(s_i) = \max_{s_{-i} \in S_{-i}} \pi_i(s_i, s_{-i}).$$

While one can imagine a very large space of possible f s, we consider a one-dimensional family that interpolates smoothly between min and max, with mean in the center. We construct such a family with following expression

$$v^\gamma(s_i) = \sum_{s_{-i} \in S_{-i}} \pi_i(s_i, s_{-i}) \cdot \frac{\exp[\gamma \cdot \pi_i(s_i, s_{-i})]}{\sum_{s \in S_{-i}} \exp[\gamma \cdot \pi_i(s_i, s)]}$$

which approaches $v^{\min}(s_i)$ as $\gamma \rightarrow -\infty$, $v^{\max}(s_i)$ as $\gamma \rightarrow \infty$, and $v^{\text{mean}}(s_i)$ when $\gamma = 0$. Intuitively, we can understand this expression as computing an expectation of the payoff for s_i under different degrees of optimism about the other player's choice of s_{-i} . In the example game above (Figure 1), the heuristic will assign highest value to s_1 (the top row) when γ is large and positive, s_2 when γ is large and negative, and s_3 when $\gamma = 0$. Notice that if $\gamma \neq 0$, the values associated with the different strategies do not necessarily correspond to a consistent belief about the other player's action. For example, if γ is positive, the highest payoff in each row will be over-weighted, but these might correspond to different columns in each row; in the example game (Figure 1) column 3 would be over-weighted when evaluating row 1 but down-weighted when evaluating rows 2 and 3. Although this internally inconsistent weighting may appear irrational, this extra degree of freedom can increase the expected payoff in a given environment without necessarily being more cognitively taxing.

Given a row-value function v , the most obvious way to select an action would be to select $\arg \max_{s_i} v$. However, exactly maximizing even a simple function may be challenging for an analog computer such as the human brain. Thus, we assume that the computation of v is subject to noise, but that this noise can be reduced through cognitive effort, which we operationalize as a single scalar φ . In particular, following Stahl and Wilson (1994), we assume that the noise is Gumbel-distributed and thus recover a multinomial logit model with the probability that player i plays strategy s_i being

$$h_{row}^{s_i}(G) = \frac{\exp[\varphi \cdot v(s_i)]}{\sum_{k \in S_i} \exp[\varphi \cdot v(k)]}.$$

Naturally, the cost of a row heuristic is a function of the cognitive effort. Specifically, we assume that the cost is proportional to effort,

$$c(h_{row}) = \varphi \cdot C_{row},$$

where $C_{row} > 0$ is a free parameter of the cost function.

3.2 Cell Heuristics

An individual might not necessarily consider all aspects connected to a strategy, but find a good “cell”, meaning payoff pair $(\pi_1(s_1, s_2), \pi_2(s_1, s_2))$. In particular, previous research has proposed that people sometimes adopt a *team view*, looking for outcomes that are good for both players, and choosing actions under the (perhaps implicit) assumption that the other player will try to achieve this mutually beneficial outcome as well (Sugden, 2003; Bacharach, 2006). Alternatively, people may engage in *virtual bargaining*, selecting the outcome that would be agreed upon if they could negotiate with the other player (Misyak and Chater, 2014). Importantly, these approaches share the assumption that people reason directly about outcomes (rather than actions) and that there is some amount of assumed cooperation.

We refer to heuristics that reason directly about outcomes as *cell heuristics*. Based on preliminary analyses, we identified one specific form of cell heuristic that participants appear to use frequently: This *jointmax* heuristic identifies the outcome that is most desirable for both players, formally

$$v^{\text{jointmax}}(s_i, s_{-i}) = \min \{ \pi_i(s_i, s_{-i}), \pi_{-i}(s_i, s_{-i}) \}$$

and the probability of playing a given strategy, with cognitive effort φ is given by

$$h_i^{s_i}(G) = \sum_{s_{-i} \in S_{-i}} \frac{\exp [\varphi \cdot v^{\text{jointmax}}(s_i, s_{-i})]}{\sum_{(k_i, k_{-i}) \in S_i \times S_{-i}} \exp [\varphi \cdot v^{\text{jointmax}}(k_i, k_{-i})]}.$$

In the example game (Figure 1), the jointmax heuristic would assign the highest probability to row **1** because the cell (**1, 3**) with payoffs (8, 8) has the highest minimum payoff.

The cognitive cost is again proportional to the accuracy, so

$$c(h_{\text{cell}}) = \varphi \cdot C_{\text{cell}},$$

where $C_{\text{cell}} > 0$ is a free parameter of the cost function.

3.3 Simulation Heuristics - Higher level reasoning

The row and cell heuristics don't construct explicit beliefs about how the other player will behave.⁴ Belief formation and best response has formed the basis of many previous models of initial play, and might very well be a sensible thing to do. We consider such a decision-making process as one possible heuristic people might use.

If a row player uses a simulation heuristic, she first considers the game from the perspective of the column player, applying some heuristic that generates a distribution of likely play. She then plays a quantal best response to that distribution. Let G' denote the transposed game, i.e., the game from the column player's perspective. Let h_{col} be the heuristic the row player use to estimate the column players behavior, then $h_{\text{sim}}(G)$ is given by

$$h_{\text{row}}^{s_r} = \frac{\exp [\varphi \cdot (\sum_{s_c \in S_c} \pi_r(s_r, s_c) \cdot h_{\text{col}}^{s_c}(G'))]}{\sum_{s_r \in S_r} \exp [\varphi \cdot (\sum_{s_c \in S_c} \pi_r(s_r, s_c) \cdot h_{\text{col}}^{s_c}(G'))]}$$

where φ is the cognitive effort parameter. A simulation heuristic is thus defined by a combination of a heuristic and a effort parameter (h_{col}, φ).

⁴They might do so implicitly, however. For example, a row heuristic that assigns a higher weight to high payoffs works well only if the other player is more likely to play those columns. Ignoring the low payoffs might correspond to an implicit belief that the other player will not play those columns.

The cognitive cost for a simulation heuristic is calculated by first calculating the cognitive cost associated with the heuristic used for the column players behavior, then a constant cost for updating the payoff matrix using that belief (C_{mul}), and one additional cost that is proportional to the cognitive effort parameter in the last step, as if it was a row heuristic,

$$c(h_{sim}) = c(h_{col}) + C_{mul} + C_{row} \cdot \varphi.$$

Notice that once the beliefs have been formed, and the beliefs have been incorporated, the last cost for taking a decision is based on C_{row} since this process is the same as averaging over the rows as for a row-heuristic.

3.4 Metaheuristic

We don't expect a person to use the same heuristic in all games. Instead, they may have a set of heuristics, choosing which one to use in each situation based on an estimate of the expected value of the candidate heuristics. We model such a process as a higher-order heuristic that selects among the first-order heuristics described above. We call this heuristic-selecting heuristic a metaheuristic.

Rather than explicitly modeling the process by which players select among the candidate heuristics, for example, using the approach in Lieder and Griffiths (2015), we use a reduced-form model based on the rational inattention model of Matějka and McKay (2015). This functional form captures the three key properties a metaheuristic should have: (1) there is a prior weight on each heuristic (2) a heuristic will be used more on games in which it is likely to perform well, and (3) the adjustment from the prior based on expected value is incomplete and costly.

Assume that an individual is choosing between n heuristics $H = \{h^1, h^2, \dots, h^N\}$. Then the probability of using heuristic h^n when playing game G is given by

$$\begin{aligned} \mathbb{P}[\{\text{use } h^n \text{ in } G\}] &= \frac{\exp[(a_n + V(h^n, \mathcal{E}, G))/\lambda]}{\sum_{j=1}^N \exp[(a_j + V(h^j, \mathcal{E}, G))/\lambda]} \\ &= \frac{p_n \exp[V(h^n, \mathcal{E}, G)/\lambda]}{\sum_{j=1}^N p_j \exp[V(h^j, \mathcal{E}, G)/\lambda]} \end{aligned} \quad (2)$$

where λ_i is an adjustment cost parameter and the a_n are weights that give the prior probability of using the different heuristics, $p_n = \frac{\exp(a_n/\lambda_i)}{\sum_{j=1}^N \exp(a_j/\lambda_i)}$.

3.4.1 The individual's optimization problem

A metaheuristic is defined by a tuple $m = \langle H, P \rangle$ where $H_i = \{h^1, h^2, \dots, h^N\}$ is a finite set of consideration heuristics, and $P = \{p^1, p^2, \dots, p^N\}$ a prior over those

heuristics. We can write down the performance of a metaheuristic in an environment \mathcal{E} , analogously to (1) for heuristics, as

$$V_i^{meta}(m, \mathcal{E}) = \sum_{G \in \mathcal{G}} \sum_{h \in H} V_i(h^n, \mathcal{E}, G) \cdot \frac{p_n \exp [V(h^n, \mathcal{E}, G)/\lambda]}{\sum_{j=1}^N p_j \exp [(V_i(h^j, \mathcal{E}, G))/\lambda]} \cdot P(G) \quad (3)$$

The optimization problem faced by the individual, subject to the adjustment cost λ , is then to maximize (3), i.e., to choose the optimal consideration set and corresponding priors,

$$m^* = \operatorname{argmax}_{H \in \mathcal{P}_{fin}(\mathcal{H})} \operatorname{argmax}_{P \in \Delta(H)} V^{meta}(\langle H, P \rangle, \mathcal{E})$$

where $\mathcal{P}_{fin}(\mathcal{H})$ is the set of all finite subsets of all possible heuristics. In practice, this is not a solvable problem when the set of possible heuristics, \mathcal{H} , is infinite. Even with a finite set of heuristics, the size of the power set will grow very quickly. We will, therefore, assume a small set of heuristics, and jointly find optimal parameters of those heuristics and priors P .

4 Experiment

Our overarching hypothesis is that individuals choose actions in one-shot games using heuristics that optimally trade off between expected payoff and cognitive cost. It is unlikely, however, that people compute these expected payoffs each time they need to make a decision. Instead, we hypothesize that people *learn* to use heuristics that are generally adaptive in their environment. This results in a critical prediction: the action a player takes on a given game will depend not only on the nature of that particular game, but also on the other games she has previously played. We tested this prediction in a large, online experiment in which participants played one-shot normal form games.

4.1 Methods

We recruited 600 participants on Amazon Mechanical Turk using the oTree platform (Chen, Schonger and Wickens, 2016). Each participant was assigned to one of 20 populations of 30 participants each. They then played 50 unique one-shot normal form games, in each period randomly matched to a new participant in their population.

Each population was assigned to one of two experimental treatments, which determined the distribution of games that were played. Specifically, we manipulated the correlation between the row and column players payoffs in each cell. In the *common interest* treatment, the payoffs were positively correlated, such that a cell with a high payoff for one player was likely to have a high payoff for the other player as well. In contrast,

in the *competing interest* treatment, the payoffs were negatively correlated, such that a cell with a high payoff for one player was likely to have a low payoff for the other. Concretely, the payoffs in each cell were sampled from a bivariate Normal distribution truncated to the range $[0, 9]$ and discretized such that all payoffs were single-digit non-negative integers.⁵ Examples of each type of *treatment game* are shown in Tables 1 and 2.

For each population, we sampled 46 treatment games, each participant playing every game once. The remaining four games were *comparison games*, treatment-independent games that we used to compare behavior in the two treatments when playing the same game. The comparison games were played in rounds 31, 38, 42, and 49. We placed them all later in the experiment so that the participants would have time to adjust to the treatment environment first, leaving gaps to minimize the chance that participants would notice that these games were different from the others they had played.

5, 6	6, 4	5, 3
9, 4	5, 5	6, 7
2, 0	0, 1	6, 4

Common interest example 1

3, 4	5, 5	9, 7
4, 2	5, 7	5, 7
2, 4	2, 1	2, 3

Common interest example 2

9, 7	5, 9	7, 8
6, 7	9, 9	4, 6
6, 4	3, 1	6, 2

Common interest example 3

1, 4	5, 3	7, 4
3, 5	4, 2	7, 5
3, 8	3, 6	5, 3

Common interest example 4

Table 1: Four games from the common interest treatment.

5, 5	6, 2	5, 3
5, 3	1, 8	8, 4
3, 6	7, 4	4, 6

Competing interest example 1

2, 4	4, 4	4, 6
1, 7	2, 6	9, 1
7, 1	4, 8	8, 6

Competing interest example 2

4, 5	1, 5	7, 1
2, 7	8, 5	5, 7
2, 6	8, 3	3, 9

Competing interest example 3

8, 0	4, 1	3, 8
4, 7	2, 7	2, 7
3, 5	3, 9	7, 5

Competing interest example 4

Table 2: Four games from the competing interest treatment.

⁵ The normal distribution is given by $N((5, 5), \Sigma)$ with $\Sigma = 5 \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$ where $\rho = 0.9$ for the common interest treatment and $\rho = -0.9$ for the competing interest treatment.

4.1.1 The Comparison Games

We selected comparison games that we expected to elicit dramatically different distributions of play in the two treatments. In all four games, there is a tension between choosing a row with an efficient outcome or a row that gives a high guaranteed pay off. For two of the games, the efficient outcome was also a Nash Equilibrium (NE), and for the other two games, the efficient outcome was not a NE.

First Comparison Game

8, 8	2, 6	0, 5
6, 2	6, 6	2, 5
5, 0	5, 2	5, 5

Comparison game 1

The first game is a weak-link game, where all the diagonal strategy profiles are Nash Equilibria, but all are not as efficient. The most efficient NE gives payoffs (8,8), but it is also possible to get 0. The least efficient equilibrium yields a payoff of (5,5), but that is also the guaranteed payoff. The equilibrium (6,6) is in between the two in terms of both risk and efficiency. The third row has the highest average payoff and is the best response to itself, so any standard level-k model would predict (5,5) being played.

Second Comparison Game

8, 8	2, 9	1, 0
9, 2	3, 3	1, 1
0, 1	1, 1	1, 1

Comparison game 2

The second comparison game is a normal prisoner's dilemma game, with an added dominated and inefficient strategy. In this game, strategy 2 is dominating the other strategies, but we still expect the common interest treatment to play strategy 1 more often since it is usually a good heuristic for them to look for the common interest.

Third Comparison Game

4, 4	4, 6	5, 0
6, 4	3, 3	5, 1
0, 5	1, 5	9, 9

Comparison game 3

The third game is game with two minimal CURB sets, $\{1, 2\}$ and $\{3\}$, where the NE inside $\{1, 2\}$ is mixed. This game is constructed so that the row averages is much higher for strategy 1 and 2 compared to 3, meaning that any level-k heuristic ends up there, while the NE yielding (9, 9) is much more efficient. So, there is a strong tension between efficiency and guaranteed payoff.

Fourth Comparison Game

4, 4	9, 1	1, 3
1, 9	8, 8	1, 8
3, 1	8, 1	3, 3

Comparison game 4

In this game, the risky efficient outcome (8, 8) is not a NE. A level-k player of any level higher than 0 would play strategy **3**.

4.2 Model fitting

The metaheuristics described above have several free parameters that control their behavior. We consider two methods for estimating these parameters. The first and more standard approach is to simply fit the parameters to the human data, optimizing the likelihood of the observed play distributions. The second approach is based on the assumption that people *rationally* adapt their heuristics to their environment. Instead of fitting the parameters of the actual heuristics, the parameters are optimized to give the highest payoff on the training games, and the cognitive costs parameters are jointly estimated for both populations to give the highest fit of these optimized heuristics on the training games. We will call these two different estimation approaches *fitting* and *optimizing*. TODO: fully describe each method, two subsections

We take an out-of-sample prediction approach to model comparison, where each data set is divided into a training and a test set and the models are estimated on the training data and evaluated on the the test data. The train data consisted of the

first 30 games of each session, and the other 16 treatment games were the test data. This split of the data was preregistered. When we estimate the optimal models, the cognitive cost parameters are chosen so that the optimized model best fit the training data. The measure of the fit we use is the negative log-likelihood (or equivalently the cross-entropy). If y is the observed distribution of play for the distribution of play for some role in some game, and x is the predicted distribution of play for some model, the negative log-likelihood (NLL) is defined

$$\text{NLL}(x, y) = - \sum_s y_s \cdot \log(x_s).$$

We then estimate the models by finding the parameters that minimize the total NLL over all the training data for some treatment.

The estimated meta heuristics consists of three heuristics, a jointmax cell heuristic, a row heuristic, and a simulation heuristic where a row heuristic models the other player’s behavior.

4.3 Results

We organize our results based on our four pre-registered hypotheses. The first two are model-free and concern behavior in the comparison games. The latter two are model-based and concern behavior in the treatment games.

4.3.1 Model-free analysis of comparison games

Our first hypothesis was that the treatment would have an effect on behavior in the comparison games.

Hypothesis 1. *The distribution of play in the four comparison games will be different in the two treatment populations.*

This hypothesis follows from the assumption that people learn to use heuristics that are adaptive in their treatment, and that different heuristics are adaptive in the two treatments. Figure 2 visually confirms this prediction and Table 3 confirms that these differences are statistically significant (χ^2 -tests, as preregistered).

	Frequencies			χ^2	p-value
	1	2	3		
Comparison Game 1				98.39	$2.2 \cdot 10^{-16}$
Common interest	193	53	54		
Competeting interest	75	82	143		
Comparison Game 2				22.08	$1.6 \cdot 10^{-5}$
Common interest	160	139	1		
Competeting interest	103	195	2		
Comparison Game 3				61.75	$3.9 \cdot 10^{-14}$
Common interest	40	73	187		
Competeting interest	106	97	97		
Comparison Game 4				91.36	$2.2 \cdot 10^{-16}$
Common interest	78	173	49		
Competeting interest	115	62	123		

Table 3: χ^2 tests for each comparison games. All of the them significant at the preregistered 0.05 level.




Figure 2: Distribution of plays in the four comparison games. Each panel shows the joint and marginal distributions of row/column plays in a single game. The cells are annotated with the each player's payoffs for the given outcome.

Inspecting Figure 2, we see that the distribution of play is not just different in the two groups; it is different in a systematic way. In particular players in the common interest treatment tend to coordinate on the efficient outcome, even in games 2 and 4, where the efficient outcome is not a Nash Equilibrium. This motivates our second hypothesis,

Hypothesis 2. *The average payoff in the four comparison games will be higher in the common interest treatment than in the competing interest treatment.*

Since the comparison games were chosen to exhibit a tension between the efficient outcome and a high guaranteed payoff, we expected that the common interest population would be better at coordinating on the efficient outcome. In our metaheuristic model, this prediction results from the fact that the jointmax heuristic generally performs quite well in the common interest games, and thus the cognitive cost of checking for each game whether another heuristic performs better generally outweighs the potential gains. Coordinating on the efficient outcomes then leads to a higher average payoff. Table 4 confirms this prediction. The common interest population had a higher average payoff in all four comparison games and the difference is significant in each case (at the pre-registered level of $p < .05$).

	Treatment average payoff			
	Common interest	Competing interest	t-value	p-value
Comparison game 1	5.09	3.64	6.851	$2 \cdot 10^{-11}$
Comparison game 2	5.52	4.04	6.28	$6.7 \cdot 10^{-10}$
Comparison game 3	5.00	4.31	2.86	0.0044
Comparison game 4	5.19	3.42	7.21	$1.9 \cdot 10^{-12}$

Table 4: Two-sided t-tests for the difference in average payoff between the two treatments in the comparison games.

4.3.2 Model-based analysis of treatment games

Next we consider our two model-based hypotheses regarding the ability of the metaheuristic model to capture the difference in strategies used in the two treatments. The first hypothesis is simply that the behavior will be different and that the model will capture some aspect of that difference.

Hypothesis 3. *The test data from each treatment should be better explained by the heuristics that are estimated on that treatment’s training set than by the heuristics estimated on the other treatment’s training set. This should hold for both estimation methods (fitting and optimizing).*

Figure 3 shows the negative log likelihood of held out test data in each treatment according to heuristic models fit to each treatment. For both the fitted and optimized

heuristics, negative log likelihood is lower (better) when training and testing on the same set (visually indicated by the outward-decreasing slope).

An even more striking (and unexpected) result is that the optimized heuristics achieved nearly the same predictive performance as the fitted heuristics. That is, a model with one set of cognitive cost parameters that applies for both treatments (with the heuristic parameters set to optimize the resultant payoff-cost tradeoff) explains participant data almost as well as the fully-parameterized model, in which the heuristic parameters are fit directly and separately to behavior in each treatment.




Figure 3: Model performance. Each panel shows the likelihood of the held-out test data for one treatment (competing vs common interest). Models are fit (blue) or optimized (orange) to either the competing (left) or common (right). Each bar represents one heuristic. The color indicates which condition the heuristic was fit to or optimized for. Fitted heuristics are optimized for likelihood on the training set of the condition corresponding to the bar’s color. Optimal heuristics are optimized for the expected payoff on the same training games in the given condition, using the empirical distribution of response plays.

Our final model-based hypothesis provides an additional test that the heuristics participants use are adapted to their treatment environment:

Hypothesis 4. *The fitted heuristics estimated for a given treatment should achieve higher expected payoffs on the test games for that treatment than should the heuristics estimated for the other treatment.*

The logic for this hypothesis is that even if we do not assume that participants use optimal heuristics, we should still see that the heuristics that best describe participant behavior in each treatment achieve higher payoffs in that treatment. The top two rows

of Table 5 confirm this prediction, although the size of the difference is small [TODO: T TEST].

Model	Estimation	Estimation data	Payoff in Common	Payoff in Competing
Metaheuristic	Fit	Competing	6.51	5.37
Metaheuristic	Fit	Common	6.69	5.33
Metaheuristic	Optimize	Competing	6.49	5.51
Metaheuristic	Optimize	Common	6.69	5.23
Random			5.38	4.45
Actual			6.74	5.43
Maximum			7.17	5.95

Table 5: The expected average payoff from using the different heuristics on the different test games. For comparison, the maximum payoff, the actual obtained average payoff, and the expected payoff from playing a random strategy are shown as well.

4.3.3 Estimated Metaheuristics

Inspecting the heuristic parameters, we find that the model captures the difference in what behavior is adaptive in each treatment in an intuitive manner. Focusing first on the fitted heuristics (Table 6), we highlight three interesting differences between the treatments. First, jointmax is used more than twice as often in the common interest condition; this aligns with the intuition that looking for an outcome that is good for both players is a better strategy when interests are usually aligned. Second, the optimism parameter of the row heuristic, γ , is lower in the competing interest treatment; this makes sense because the better outcomes for you are less likely to be realized when your partner has opposing interests. Third, simulation is used more than twice as often in the competing interest treatment; this suggests that careful reasoning about other agents is more important in competitive environments—or conversely, simple heuristics like jointmax are more successful in cooperative environments.

Fitted Metaheuristics				
		Jointmax	Row Heuristic	Sim
		φ	γ, φ	$(\gamma, \varphi), \varphi$
Common interests	Params	2.47	-0.26, 1.79	(-0.26, 0.05), 0.97
	Share	35 %	48 %	17 %
Competing interests	Params	1.47	-1.34, 2.66	(1.45, 0.51), 0.96
	Share	15 %	47 %	37 %

Table 6: The estimated meta heuristics for the two treatments.

Optimal Metaheuristics				
		Jointmax	Row Heuristic	Sim
		φ	γ, φ	$(\gamma, \varphi), \varphi$
Common interests	Params	1.23	1.2, 1.43	(0.64, 0.53), 1.26
	Share	45 %	53 %	2 %
Competing interests	Params	0.98	-1.52, 1.17	(-1.11, 0.67), 1.13
	Share	0 %	70 %	30 %

Table 7: The optimal meta heuristics for the two treatments.

Table 6 shows the parameters for the optimized heuristics. All three of the patterns we highlight above are replicated, although they are all stronger. This perhaps suggests that participants’ adaptation to the experimental environment was only partial, which is not surprising given the relative weight of 30-46 games in an experiment versus years of experience interacting with people in the real world.

5 Deep Heuristics

A drawback of using the explicitly formulated heuristics that we use to build up the metaheuristics is that it is dependent on the decisions made by the researchers. The space of possible heuristics is extremely large, and the ultimately somewhat arbitrary decisions the researcher makes about which heuristics to include might affect the ultimate results. To minimize the risk of our conclusions being driven by such decisions, we also take a very different route to modeling heuristics. While not lending itself to as clear interpretability as the metaheuristics, it includes a much larger set of possible heuristics.

To get a general model of possible heuristics, we use a neural network architecture close to the one developed in Hartford, Wright and Leyton-Brown (2016), with some adjustments to improve interpretability and easier modeling of cognitive costs. The idea is to let every element of the input and hidden layers be a matrix, instead of a single value, of the same size as the game. Each cell in those matrices is then treated in the same way. This ensures that the deep heuristic is invariant to relabeling of strategies, as should be expected from any decision rule for normal-form games.

Higher-level reasoning is incorporated by first having two separated neural networks, representing the heuristic for the row player and the column player separately, and then possibly take into account the thus formed beliefs about the column player’s behavior in the action response layers. The different action response layers are then mixed into a response. A heuristic that does not explicitly form beliefs about the other player’s behavior would let $AR^R(0)$ be the output, a person who only applies a heuristic to estimate the opponent’s behavior and then best responds to it would only use $AR^R(1)$, etc.

5.1 Feature Layers

The hidden layers are updated according to

$$H_{l,k}^R = \phi_l \left(\sum_j w_{l,k,j}^R H_{l-1,j}^R + b_{l,k}^R \right) \quad H_{l,k}^R \in \mathbb{R}^{m_R \times m_C}$$

and similarly for H^C . For the first hidden layer $H_{0,i}^R = H_{0,i}^C = F_i$, so the two disjoint parts start with the same feature matrices, but then have different weights.

The feature matrices consist of matrices where each cell contains information associated with the row or column of one payoff matrix. The payoff matrices for the row and column players are denoted U^R and U^C , respectively. More specifically, we calculate the maximum, minimum, and mean of each row and column for both payoff matrices. Furthermore, F_1 and F_2 are the payoff matrices as they are, and lastly, we have a feature matrix where each value is the minimum payoff that either one of the players receives from the strategy profile. Below follow three examples of such feature matrices.

$$\begin{pmatrix} \max_i U_{i,1}^R & \max_i U_{i,2}^R & \max_i U_{i,3}^R \\ \max_i U_{i,1}^R & \max_i U_{i,2}^R & \max_i U_{i,3}^R \\ \max_i U_{i,1}^R & \max_i U_{i,2}^R & \max_i U_{i,3}^R \end{pmatrix}, \quad \begin{pmatrix} \max_j U_{1,j}^R & \max_j U_{1,j}^R & \max_j U_{1,j}^R \\ \max_j U_{2,j}^R & \max_j U_{2,j}^R & \max_j U_{2,j}^R \\ \max_j U_{3,j}^R & \max_j U_{3,j}^R & \max_j U_{3,j}^R \end{pmatrix}$$

$$\begin{pmatrix} \min_{R,C} \{U_{1,1}^R, U_{1,1}^C\} & \min_{R,C} \{U_{1,2}^R, U_{1,2}^C\} & \min_{R,C} \{U_{1,3}^R, U_{1,3}^C\} \\ \min_{R,C} \{U_{2,1}^R, U_{2,1}^C\} & \min_{R,C} \{U_{2,2}^R, U_{2,2}^C\} & \min_{R,C} \{U_{2,3}^R, U_{2,3}^C\} \\ \min_{R,C} \{U_{3,1}^R, U_{3,1}^C\} & \min_{R,C} \{U_{3,2}^R, U_{3,2}^C\} & \min_{R,C} \{U_{3,3}^R, U_{3,3}^C\} \end{pmatrix}$$

Figure 4: Examples of input feature matrices.

5.2 Softmax and Action Response Layers

After the last feature layer, a play distribution is calculated from each feature matrix in the last layer. This is done by first summing over the rows (columns) and then taking a softmax over the sums. The first action response layer is then given by a weighted average of those different distributions. So for example, the distribution $S_1^R \in \Delta^{m_R}$ is give by

$$S_1^R = \text{softmax} \left(\sum_i (H_{2,1}^R)_{1,i}, \sum_i (H_{2,1}^R)_{2,i}, \dots, \sum_i (H_{2,1}^R)_{m_R,i} \right)$$

while the sums for the column player is taken over the columns, so

$$S_1^C = \text{softmax} \left(\sum_j (H_{2,1}^C)_{j,1}, \sum_j (H_{2,1}^C)_{j,2}, \dots, \sum_j (H_{2,1}^C)_{j,m_C} \right).$$

The first action response distribution is then $\text{AR}^R(0) = \sum_l^{k_2} w_l^R S_l^R$ for $w^R \in \Delta^{k_2}$, and similarly for the column player.

The $\text{AR}^R(0)$ corresponds to a level-0 heuristic, a heuristic where the column player's behavior isn't explicitly modeled and taken into account. To do this, we move to Action Response layer 1, and use $\text{AR}^C(0)$ as a prediction for the behavior of the opposing player. Once the beliefs for the play of the column player are formed, the $\text{AR}^R(1)$ calculates the expected value from each action, conditioned on that expected play, and takes a softmax over those payoffs.

$$\text{AR}^R(1) = \text{softmax} \left(\lambda \sum_j U_{1,j}^R \cdot \text{AR}^C(0)_j, \dots, \lambda \sum_j U_{m_R,j}^R \cdot \text{AR}^C(0)_j \right)$$

Higher-level action response layers form its beliefs about the other player by taking a weighted average of the lower layers and respond in the same fashion.

5.3 Output layer

The output layer takes a weighted average of the row player's action response layers. This is the final predicted distribution of play for the row player.

5.4 Cognitive costs

When the deep heuristic is optimized with respect to received payoff, the cognitive cost comes from two features of the network. Firstly, there is an assumed fixed cost associated with simulating, so the higher proportion is given to $\text{AR}^R(1)$, the higher that cost. Secondly, it is assumed that more exact predictions are more cognitively costly. The second cognitive cost is thus proportional to the reciprocal of the entropy of the resulting prediction.

5.5 Results

By applying the same estimation method to the deep heuristics, we can test if hypotheses 3 and 4 also hold for a completely different specification of the space of heuristics and cognitive costs. In Figure 5, we see that Hypothesis 3 holds for this specification as well. We also see that the predictive performance of the optimal heuristic is close to the fitted heuristic, given optimized cognitive costs.




Figure 5: Model performance for the deep heuristic. Each panel shows the likelihood of the held-out test data for one treatment (competing vs common interest). Models are fit (blue) or optimized (orange) to either the competing (left) or common (right) treatment. Each bar represents one heuristic. The color indicates which condition the heuristic was fit to or optimized for. Fitted heuristics are optimized for likelihood on the training set of the condition corresponding to the bar’s color. Optimal heuristics are optimized for the expected payoff on the same training games in the given condition, using the empirical distribution of response plays.

We can also test Hypothesis 4 in the same way at looking at the expected payoff from the two different deep heuristics fitted to the behavior of the populations in the two different treatments, and see that this also holds for the deep heuristics.

Model	Estimation	Estimation data	Payoff in Common	Payoff in Competing
Deep heuristic	Fit	Competing	6.41	5.37
Deep heuristic	Fit	Common	6.66	5.06
Deep heuristic	Optimize	Competing	6.35	5.44
Deep heuristic	Optimize	Common	6.71	5.07
Random			5.38	4.45
Actual			6.74	5.43
Maximum			7.17	5.95

Table 8: The expected average payoff from using the different deep heuristics on the different test games. For comparison, the maximum expected payoff, the actual obtained average payoff, and the expected payoff from playing a random strategy is shown as well.

6 Alternative Models

In the previous sections we have seen that rational use of heuristics can explain and predict behavior in one-shot games, and that this is a result we can reproduce with two very different specifications of the space of heuristics. To further show the appropriateness of the chosen spaces of heuristics, and the strength of the predictions, we consider two alternative models of behavior: quantal cognitive hierarchy and prosocial preferences.

Quantal Cognitive Hierarchy. In previous comparisons, variations of cognitive hierarchy models are usually the best performing, Camerer, Ho and Chong (2004); Wright and Leyton-Brown (2017). In such a model, we consider agents of different cognitive levels. In the quantal cognitive hierarchy model we consider here, a level-0 agent plays the uniformly random strategy, playing each action with an equal probability. Level-1 plays a quantal best response to a level-0, and a level-2 player best responds to a combination of level-0 and level-1. In total this model has 4 parameters, the share of level-0 and level-1 players (and thus also the share of level-2), the sensitivity λ_1 of level-1 players and the sensitivity λ_2 of level-2 players.

Prosocial preferences. In the comparison games where we see that difference in behavior could potentially be explained by simply having correct and differing beliefs and potentially some kind of pro-social preferences. Since the populations behavior differ, one possible explanation to observed differences in behavior might simply be that the two populations “coordinate” on different norms without necessarily relying explicitly on optimal heuristics. We therefore test a model with noisy best reply to the correct beliefs with a prosocial utility function. The prosocial utility function we consider is

$$u_i(s_i, s_{-i}) = (1 - \alpha s - \beta r) \times \pi_i(s_i, s_{-i}) + (\alpha s + \beta r) \times \pi_{-i}(s_i, s_{-i})$$

where s indicates if $\pi_i(s_i, s_{-i}) < \pi_{-i}(s_i, s_{-i})$ and r indicates if $\pi_i(s_i, s_{-i}) > \pi_{-i}(s_i, s_{-i})$. In other words α determines how much player i values the payoff of player $-i$ when i gains less, and β how much player i values the payoff of player $-i$ when i gain more.

In table 9 we compare the out of sample predictive performance of these two alternative models and our two suggested specifications for the space of heuristics.

Table 9: Out of sample average negative log-likelihood for common interest test games.

Model	Average NLL
Fitted Meta Heuristic	0.596
Optimal Meta heuristic	0.597
Fitted Deep Heuristic	0.592
Optimal Deep Heuristic	0.643
Quantal Cognitive Hierarchy	0.686
Prosocial Preferences	0.717

Table 10: Out of sample average negative log-likelihood for competing interest test games.

Model	Average NLL
Fitted Meta Heuristic	0.714
Optimal Meta heuristic	0.735
Fitted Deep Heuristic	0.705
Optimal Deep Heuristic	0.770
Quantal Cognitive Hierarchy	0.736
Prosocial Preferences	0.838

As we see, the fitted versions of our heuristics outperform the alternatives in both treatments, and with a lot in the common interest games. The same is true for the versions with optimal instead of estimated heuristics on the common interest treatment games, but on the competing interest games the QCH model performs similar to the optimal Meta Heuristic, and better than the optimal Deep Heuristic. Adding the losses from the two treatments together however, the optimal versions our heuristics outperform the estimated QCH.

In conclusion, the models we propose clearly outperform alternative explanations.

7 Discussion

In the theory presented we combine two perspectives. On the one hand we assume that people use simple cognitive process, working directly on the level of the payoff information, to choose actions which are in often inconsistent with rational behavior in any give game. On the other hand, we don't assume that the specific heuristics used are predetermined or insensitive to incentives. On the contrary, we assume that the specific heuristics people use are chosen according to a resource-rational analysis such that they strike an optimal balance between expected payoffs and cognitive costs. We

have seen that by combining these two perspectives, we can get better predictions of behavior and understand the influence of the larger environment on the behavior in a given game.

This approach gives a lens through which we can understand the connection between classical rational models of behavior and more behavioral approaches. More generally, it can help us predict when we should expect behavior to coincide with rational behavior and when we might see systematic deviations from a rational benchmark. If the considered decision situation is such that there exists a simple heuristic which coincides with the optimal behavior, and this heuristic is good in the general environment, we can expect behavior to appear rational. If there is no simple heuristic that leads to optimal behavior in an environment, or if the situation differs from the general environment in such a way that the usually optimal heuristic is performing poorly, we should expect consistent deviations from the rational benchmark. The comparison games illustrate an important consequence of the optimal use of heuristics. The optimal heuristic will focus on the features of the games that are often of importance, but miss opportunities that are rare. So a person used to common interest games might miss the beneficial deviations from the efficient outcome, while persons used to a competing interest games fail to see and take advantage of a common interest.

Our findings relate to those of Peysakhovich and Rand (2016), where varying the sustainability of cooperation in an initial session of repeated prisoner’s dilemma affected how much pro-social behavior and trust was shown in later games, including one-shot prisoner’s dilemma. Our results provide a qualitative replication of this idea. In particular, we found that putting people in an environment in which pro-social heuristics (such as jointmax) perform well led them to choose pro-social actions in the comparison games, in some cases, even selecting dominated options. In contrast, putting people in an environment where pro-social actions often result in low payoffs prevented people from achieving efficient outcomes, even when they were Nash Equilibria. Consistent with our theory, the authors interpreted their findings as the result of heuristic decision making. We build on this intuitively appealing notion by specifying a formal model of heuristics in one-shot games that makes quantitative predictions. We also emphasize the influence of cognitive costs (in addition to payoffs) on the heuristics people use.

Lastly, a point relating to the larger literature. In our theory, the generalization between games happens on the level of reasoning; the individuals are not learning which actions are good, but rather how they should choose the action. This contrasts with theories where the generalization happens on the level of actions, as in Jehiel (2005) or Grimm and Mengel (2012). Furthermore, since our games are randomly generated, no such action learning should take place or alter behavior systematically.

8 Conclusion

We have proposed a theory of human behavior in one-shot normal form games based on the resource-rational use of heuristics. According to this theory, people select their actions using simple cognitive heuristics that flexibly and selectively process payoff information; the heuristics people choose to use are ones that strike an optimal tradeoff between payoffs and cognitive cost.

In a large preregistered experiment, we confirmed one of the primary qualitative predictions of the theory: people learn what heuristics are resource-rational in a given environment, and thus their recent experience affects the choices they make. In particular, we found that placing participants in environments with common (vs. competing) interests leads them to select the most efficient (or least efficient) equilibrium in a weak link game and to cooperate (or defect) in prisoner’s dilemma.

Furthermore, we found that our theory provides a strong quantitative account of our participants’ behavior, making more accurate out-of-sample predictions than both the state of the art quantal cognitive hierarchy model (Wright and Leyton-Brown, 2017) and also a prosocial preference model. Strikingly, we found that a resource-rational model, in which behavior in both treatments is predicted using a single set of fitted cost parameters (with the heuristic parameters set to optimize the resultant payoff-cost tradeoff), achieved nearly the same accuracy as the fully-parameterized model, in which the heuristic parameters are fit directly and separately to behavior in each treatment. Coupled with the overall high predictive accuracy of the model, this provides strong evidence in support of the theory that people use heuristics that optimally trade off between payoff and cognitive costs. In a followup analysis, we found similar results using an entirely different neural-network-based family of heuristics, indicating that these findings are robust to the parameterization of the heuristics.

From a wider perspective, our model speaks to a decades-long debate on the rationality of human decision making. With classical models based on optimization and utility maximization failing to capture systematic patterns in human choice behavior, recent models instead emphasize our systematic biases, suggesting that we rely on simple and error-prone heuristics to make decisions. In this paper, we hope to have offered a synthesis of these two perspectives, by treating heuristics as things that can themselves be optimized in a utility-maximization framework. We hope that this approach will be valuable in working towards a more unified understanding of economic decision making.

References

Bacharach, Michael. 2006. *Beyond individual choice: teams and frames in game theory*. Princeton University Press.

- Bardsley, Nicholas, Judith Mehta, Chris Starmer, and Robert Sugden.** 2010. "Explaining focal points: Cognitive hierarchy theory versus team reasoning." *Economic Journal*, 120(543): 40–79.
- Camerer, C. F., T.-H. Ho, and J.-K. Chong.** 2004. "A Cognitive Hierarchy Model of Games." *The Quarterly Journal of Economics*, 119(3): 861–898.
- Camerer, Colin F.** 2003. *Behavioral Game Theory: Experiments in Strategic Interaction*. Princeton University Press.
- Caplin, Andrew, and Mark Dean.** 2013. "Behavioral Implications of Rational Inattention with Shannon Entropy." *NBER Working Paper*, , (August): 1–40.
- Chen, Daniel L., Martin Schonger, and Chris Wickens.** 2016. "oTree—An Open-Source Platform for Laboratory, Online, and Field Experiments." *Journal of Behavioral and Experimental Finance*, 9: 88–97.
- Costa-Gomes, Miguel A., and Georg Weizsäcker.** 2008. "Stated Beliefs and Play in Normal-Form Games." *Review of Economic Studies*, 75(3): 729–762.
- Crawford, Vincent P, Miguel A Costa-Gomes, and Nagore Iriberri.** 2013. "Structural Models of Nonequilibrium Strategic Thinking: Theory, Evidence, and Applications." *Journal of Economic Literature*, 51(1): 5–62.
- Devetag, Giovanna, Sibilla Di Guida, and Luca Polonio.** 2016. "An eye-tracking study of feature-based choice in one-shot games." *Experimental Economics*, 19(1): 177–201.
- Ert, Eyal, and Ido Erev.** 2013. "On the descriptive value of loss aversion in decisions under risk: Six clarifications." *Judgment and Decision Making*, 8(3): 214–235.
- Fudenberg, Drew, and Annie Liang.** 2019. "Predicting and Understanding Initial Play." *American Economic Review*, 109(12): 4112–4141.
- Fudenberg, Drew, Fudenberg Drew, David K Levine, and David K Levine.** 1998. *The theory of learning in games*. Vol. 2, MIT press.
- Gershman, S. J., E. J. Horvitz, and J. B. Tenenbaum.** 2015. "Computational Rationality: A Converging Paradigm for Intelligence in Brains, Minds, and Machines." *Science*, 349(6245).
- Gigerenzer, Gerd, and Peter M Todd.** 1999. *Simple Heuristics That Make Us Smart*. Oxford University Press, USA.
- Goldstein, Daniel G., and Gerd Gigerenzer.** 2002. "Models of Ecological Rationality: The Recognition Heuristic." *Psychological review*, 109(1): 75.
- Griffiths, Thomas L, Falk Lieder, and Noah D Goodman.** 2015. "Rational Use of Cognitive Resources: Levels of Analysis between the Computational and the Algorithmic." *Topics in Cognitive Science*, 7(2): 217–229.

- Grimm, Veronika, and Friederike Mengel.** 2012. “An experiment on learning in a multiple games environment.” *Journal of Economic Theory*, 147(6): 2220–2259.
- Hartford, Jason S., James R. Wright, and Kevin Leyton-Brown.** 2016. “Deep Learning for Predicting Human Strategic Behavior.” *Advances in Neural Information Processing Systems*, , (Nips): 2424–2432.
- Heap, Shaun Hargreaves, David Rojo Arjona, and Robert Sugden.** 2014. “How Portable is Level-0 Behavior? A Test of Level-k Theory in Games with Non-Neutral Frames.” *Econometrica*, 82(3): 1133–1151.
- Hebert, Benjamin, and Michael Woodford.** 2019. “Rational Inattention When Decisions Take Time.” *Journal of Chemical Information and Modeling*, 53(9): 1689–1699.
- Imai, Taisuke, Tom A Rutter, and Colin F Camerer.** 2020. “Meta-Analysis of Present-Bias Estimation Using Convex Time Budgets*.” *The Economic Journal*, 186(2): 227–236.
- Jehiel, Philippe.** 2005. “Analogy-based expectation equilibrium.” *Journal of Economic Theory*, 123(2): 81–104.
- Kahneman, Daniel.** 2011. *Thinking, fast and slow*. Macmillan.
- Lewis, Richard L., Andrew Howes, and Satinder Singh.** 2014. “Computational Rationality: Linking Mechanism and Behavior through Bounded Utility Maximization.” *Topics in Cognitive Science*, 6(2): 279–311.
- Lieder, Falk, and Thomas L. Griffiths.** 2015. “When to use which heuristic: A rational solution to the strategy selection problem.” *Proceedings of the 37th annual conference of the cognitive science society*, 1(3): 1–6.
- Lieder, Falk, and Thomas L. Griffiths.** 2017. “Strategy Selection as Rational Metareasoning.” *Psychological Review*, 124(6): 762–794.
- Lieder, Falk, and Thomas L. Griffiths.** 2019. “Resource-Rational Analysis: Understanding Human Cognition as the Optimal Use of Limited Computational Resources.” *Behavioral and Brain Sciences*.
- Lieder, Falk, Paul M Krueger, and Tom Griffiths.** 2017. “An automatic method for discovering rational heuristics for risky choice.”
- Matějka, Filip, and Alisdair McKay.** 2015. “Rational Inattention to Discrete Choices: A New Foundation for the Multinomial Logit Model.” *American Economic Review*, 105(1): 272–298.
- Mengel, Friederike, and Emanuela Sciubba.** 2014. “Extrapolation and structural similarity in games.” *Economics Letters*, 125(3): 381–385.

- Misyak, Jennifer B., and Nick Chater.** 2014. "Virtual Bargaining: A Theory of Social Decision-Making." *Philosophical Transactions of the Royal Society B: Biological Sciences*, 369(1655).
- Nagel, Rosemarie.** 1995. "Unraveling the Guessing Game." *American Economic Review*, 85(5): 1313–1326.
- Peysakhovich, Alexander, and David G. Rand.** 2016. "Habits of virtue: Creating norms of cooperation and defection in the laboratory." *Management Science*, 62(3).
- Polonio, Luca, Sibilla Di Guida, and Giorgio Coricelli.** 2015. "Strategic sophistication and attention in games: An eye-tracking study." *Games and Economic Behavior*, 94: 80–96.
- Savage, Leonard J.** 1954. *The Foundations of Statistics. The Foundations of Statistics*, Oxford, England: John Wiley & Sons.
- Simon, Herbert A.** 1976. "From substantive to procedural rationality." *25 Years of Economic Theory*, 65–86.
- Sims, C. A.** 1998. "Stickiness." *Carnegie-Rochester Conference Series on Public Policy*, 49: 317–356.
- Smith, Vernon L.** 2003. "Constructivist and Ecological Rationality in Economics." *American economic review*, 93(3): 465–508.
- Stahl, Dale O., and Paul W. Wilson.** 1994. "Experimental evidence on players' models of other players." *Journal of Economic Behavior and Organization*, 25(3): 309–327.
- Stahl, Dale O., and Paul W. Wilson.** 1995. "On players' models of other players: Theory and experimental evidence."
- Steiner, Jakub, Colin Stewart, and Filip Matějka.** 2017. "Rational Inattention Dynamics: Inertia and Delay in Decision-Making." *Econometrica*, 85(2): 521–553.
- Stewart, Neil, Simon Gächter, Takao Noguchi, and Timothy L Mullett.** 2016. "Eye Movements in Strategic Choice." 156(October 2015): 137–156.
- Sugden, Robert.** 2003. "The logic of team reasoning." *Philosophical explorations*, 6(3): 165–181.
- Todd, Peter M., and Gerd Ed Gigerenzer.** 2012. *Ecological Rationality: Intelligence in the World*. Oxford University Press.
- Tunçel, Tuba, and James K. Hammitt.** 2014. "A new meta-analysis on the WTP/WTB disparity." *Journal of Environmental Economics and Management*, 68(1): 175–187.

Wright, James R., and Kevin Leyton-Brown. 2017. “Predicting human behavior in unrepeated, simultaneous-move games.” *Games and Economic Behavior*, 106(2): 16–37.

A Instructions for the experiment

Instructions

In this HIT you will play 50 two-player games with many different real people. In each game, you will see a table like the one below. You will choose one of the three rows, and the other person will choose a column in the same way. These two decisions select one cell from the table, which determines the points you will each receive.

3 3	0 6	1 5
6 0	9 0	2 6
2 3	4 8	8 1

In each cell, there are two numbers. The first (orange) number is the number of points you get, and the second (blue) number is the number of points the other person gets. These points will determine the bonus payment you receive at the end of the HIT. For example, if you choose the third row and the other person chooses the second column, you would receive 4 points and she or he would receive 8 points, as shown below.

3 3	0 6	1 5
6 0	9 0	2 6
2 3	4 8	8 1

You will be playing against real people. For each game, you will be matched with a **new person**. To keep things moving quickly, you will sometimes be matched with a player who has already played the game in a previous round. Although your move will not affect that player's score, it will affect future players that get matched with you, just as your score is determined by the previous player's move.

Because you are playing against real people, there may be a delay after the first game while other players complete the instructions. Please be patient! It should go much faster for the remaining games. You will be compensated with an extra bonus payment for the time spent on wait pages at a rate of **\$7 an hour**.

Your bonus will be determined by the total number of points you earn in the experiment. You will get **\$1** bonus payment for each **150 points**.

One last thing. To prevent people from quickly clicking through the experiment without thinking, we enforce that you spend a minimum of 5 seconds on each game.

Before beginning to play, you must pass a quiz to demonstrate that you understand the rules. You must pass all three pages of the quiz before you can continue.

Next

Figure 6: The instructions one the first page when a participant joins the experiment.

Quiz 1 of 3

To ensure that you understand the rules, please answer the questions below. If you answer any question incorrectly, you will be brought back to the Instructions page to review.

5 8	6 6	6 6
2 3	1 7	3 7
4 2	4 4	1 7

You choose the **third** row and the other person chooses the **third** column.

What payoff do you receive?

What payoff does the other player receive?

Next

Figure 7: The participants have to complete three questions like this in a row in order to be allowed to participate in the experiment.

Round 1 of 50

3 0	7 6	2 3
4 5	5 4	5 6
7 9	3 3	4 1

Please choose a row.

Next

Figure 8: In each round, the participant chose a row by clicking on it. Once it is clicked it is highlighted and they have to click the next button to proceed.

Result




Figure 9: Once the behavior of the matched participant, either by her making a decision or by sampling from previous decisions in the game from the same population, the result is shown.

B Explaining adaptation via learning

In the main text, we assume that the participants manage to find rational heuristics without going into the details about how that is done. Here, we show that a learning model could explain this adaptation to the rational heuristics.

For simplicity, we model the learning process on the level of the population. We assume that all individuals arrive at the experiment with the same initial metaheuristic $mh(\cdot|\theta(0))$ where θ are the parameters of the metaheuristic. For each experimental session ξ , the players play a sequence of $(G_{\xi,t})_{t=1}^{50}$. Given the actual play we can evaluate the expected utility from using the metaheuristic mh with parameters θ as the value of $V(mh(\cdot, \theta), \mathcal{E}, G)$, i.e. the expected payoff minus cognitive cost for using the metaheuristic with parameters θ in game G and environment \mathcal{E} . The behavior of the other players, i.e. \mathcal{E} , is given by the actual behavior.

A straightforward learning model would be for the individuals in a session to at each timestamp updated the parameters according to what would have improved the outcome in the game. In other words,

$$\theta_{\xi}(t+1) = \theta_{\xi}(t) + \lambda \nabla_{\theta} V^{meta}(mh(\cdot, \theta), \mathcal{E}, G_{\xi,t}) \quad (4)$$

$$\theta_{\xi}(0) = \theta(0), \quad (5)$$

where λ is a learning parameter governing the speed of learning. Essentially, the learning follows a gradient decent over the parameters.

We use the costs estimated for the optimal heuristics $C_{row} = 0.41$, $C_{cell} = 0.30$, $C_{multiplication} = 0.13$, $\lambda_{mh} = 2.1$. To estimate this model we thus need a baseline heuristic, i.e. $\theta(0)$ and a learning parameter λ . To make the performance of this model comparable to that of the other models, we estimate the common starting parameters $\theta(0)$ and the common learning rate λ in order to minimize loss on the first 30 games of each session in both treatments. We then predict the last 16 games of each session in both treatments.

Table 11: Out of sample average negative log-likelihood for the learning model.

Model	Competing interest Avg NLL	Common interest Avg NLL
Fitted Meta Heuristic	0.714	0.596
Optimal Meta heuristic	0.735	0.597
Fitted Deep Heuristic	0.705	0.598
Optimal Deep Heuristic	0.770	0.643
Learning	0.723	0.608

In table 11 we see that the performance of the learning model is comparable to, but ever so slightly worse than the fitted models.

B.1 A different learning model

In the metaheuristic, there is a prior probability p_n of using each primitive heuristic h_n , which is captured in equation 2. The prior probability can be expressed as

$$p_n = \frac{\exp(a_n/\lambda_i)}{\sum_{j=1}^N \exp(a_j/\lambda_i)},$$

where a_n can be interpreted as the ex ante expected utility from using heuristic n .

In this learning model we assume that the participants start with a common initial expectation of the value of the different heuristics. As the games are played, this expectation is updated according to the actual value of the different primitive heuristics in each game. Let $a_n(1)$ be the initial expected utility from using the primitive heuristic h_n , equal for all participants in both treatments. Let the environment be determined by the actual distribution of play in each game. For a given session, we then update these attractions after each game G_t according to

$$a_n(t+1) = \rho a_n(t) + \beta \frac{V(h_n, \mathcal{E}, G_t) + V(h_n, \mathcal{E}, G'_t)}{2} \quad (6)$$

where ρ is a recency parameter decaying old information, and β determines the influence of new information. We take the average of the value of the heuristic from row players perspective G and from the column players perspective G' since each player is randomly assigned a role in each game.

To estimate this learning model we thus specify a set of primitive heuristics $H = (h_n)_{n=1}^N$, whose parameters remain fixed, their initial attractions a_n , and the two learning parameters ρ and β .

In the main specification of the metaheuristics, we consider three primitive heuristics for each treatment. With this learning model, we assume that the primitive heuristics have constant parameters and sensitivities, and leaning is only on the level of the prior probability of using these different primitive heuristics. We therefore consider a larger set of primitive heuristics so that we may still capture effects such as different row heuristics being used in the different treatments.

We consider the following 6 primitive heuristics: one common interest heuristic, one optimistic row heuristic ($\gamma = 2$), one neutral row heuristic $\gamma = 0$, one pessimistic row heuristic ($\gamma = -2$), and lastly one simulation heuristic responding to a neutral row heuristic.

To make the out of sample predictions comparable, we estimate the parameters of the learning model in order to minimize to loglikelihood on the training data, i.e., the first 30 games in each experimental session.